

# 上机实验八 & 实验九 实验报告

刘瀚文

517030910294

2018 年 11 月 20 日

# 目录

<b>I</b>	<b>实验八</b>	<b>3</b>
<b>1</b>	<b>实验准备</b>	<b>4</b>
1.1	环境配置 . . . . .	4
1.1.1	hadoop 安装 . . . . .	4
1.2	背景知识 . . . . .	5
1.2.1	Brief Introduction for Hadoop . . . . .	5
1.2.2	MapReduce Overview . . . . .	6
1.2.3	Map . . . . .	6
1.2.4	Architecture . . . . .	7
1.3	简单的自我尝试 . . . . .	7
<b>2</b>	<b>Mini Exercise</b>	<b>8</b>
2.1	Exercise 1 . . . . .	8
2.1.1	要求 . . . . .	8
2.1.2	实验过程 . . . . .	8
2.2	Exercise 2 . . . . .	11
2.2.1	要求 . . . . .	11
2.2.2	实验过程 . . . . .	11
<b>II</b>	<b>实验总结</b>	<b>13</b>
<b>III</b>	<b>实验 9</b>	<b>15</b>
<b>3</b>	<b>实验准备</b>	<b>17</b>

3.1	背景知识 . . . . .	17
3.1.1	HDFS Concepts . . . . .	17
3.1.2	Hadoop Streaming . . . . .	18
<b>4</b>	<b>实验操作</b>	<b>19</b>
4.1	exercise1 . . . . .	19
4.1.1	mapper.py . . . . .	19
4.1.2	reducer.py . . . . .	19
4.1.3	结果展示 . . . . .	21
4.2	exercise2 . . . . .	21
4.2.1	mapper.py . . . . .	21
4.2.2	结果展示 . . . . .	22
4.2.3	在第一次完成实验的基础上，我又重新写了一份代码 使用 mapper + reducer 的方法来进行操作 . . . . .	22
4.2.4	实验结果展示 . . . . .	24
<b>IV</b>	<b>实验总结</b>	<b>27</b>

## Part I

## 实验八

# Chapter 1

## 实验准备

### 1.1 环境配置

#### 1.1.1 hadoop 安装

主要的难点就是 hadoop 的配置和安装，完成了环境的配置就几乎要成功了。

#### Prerequisite

**Install Ubuntu**

**Create Hadoop User**

**Setup SSH Certification**

**Install Java and ssh-server**

**Download Hadoop 2.2.0**

**Setup Hadoop Enironment**

**Configure Hadoop**

**Format Namenode**

**Start Hadoop Service**

**Stop Services**

## **1.2 背景知识**

### **1.2.1 Brief Introduction for Hadoop**

#### **Hadoop 简介**

Formally speaking, Hadoop is an open source framework for writing and running distributed applications that process large amounts of data.

A Hadoop cluster has many parallel machines that store and process large data sets. Client computers send jobs into this computer cloud and obtain results.

#### **Hadoop 优点**

**Accessible** Hadoop runs on large clusters of commodity machines or on cloud computing services such as Amazon’ s Elastic Compute Cloud (EC2 ).

**Robust** Because it is intended to run on commodity hardware, Hadoop is architected with the assumption of frequent hardware malfunctions. It can gracefully handle most such failures.

**Scalable** Hadoop scales linearly to handle larger data by adding more nodes to the cluster.

**Simple** Hadoop allows users to quickly write efficient parallel code.

## 1.2.2 MapReduce Overview

### Characteristic

- Automatic parallelization & distribution
- Fault-tolerant
- Provides status and monitoring tools
- Clean abstraction for programmers

## 1.2.3 Map

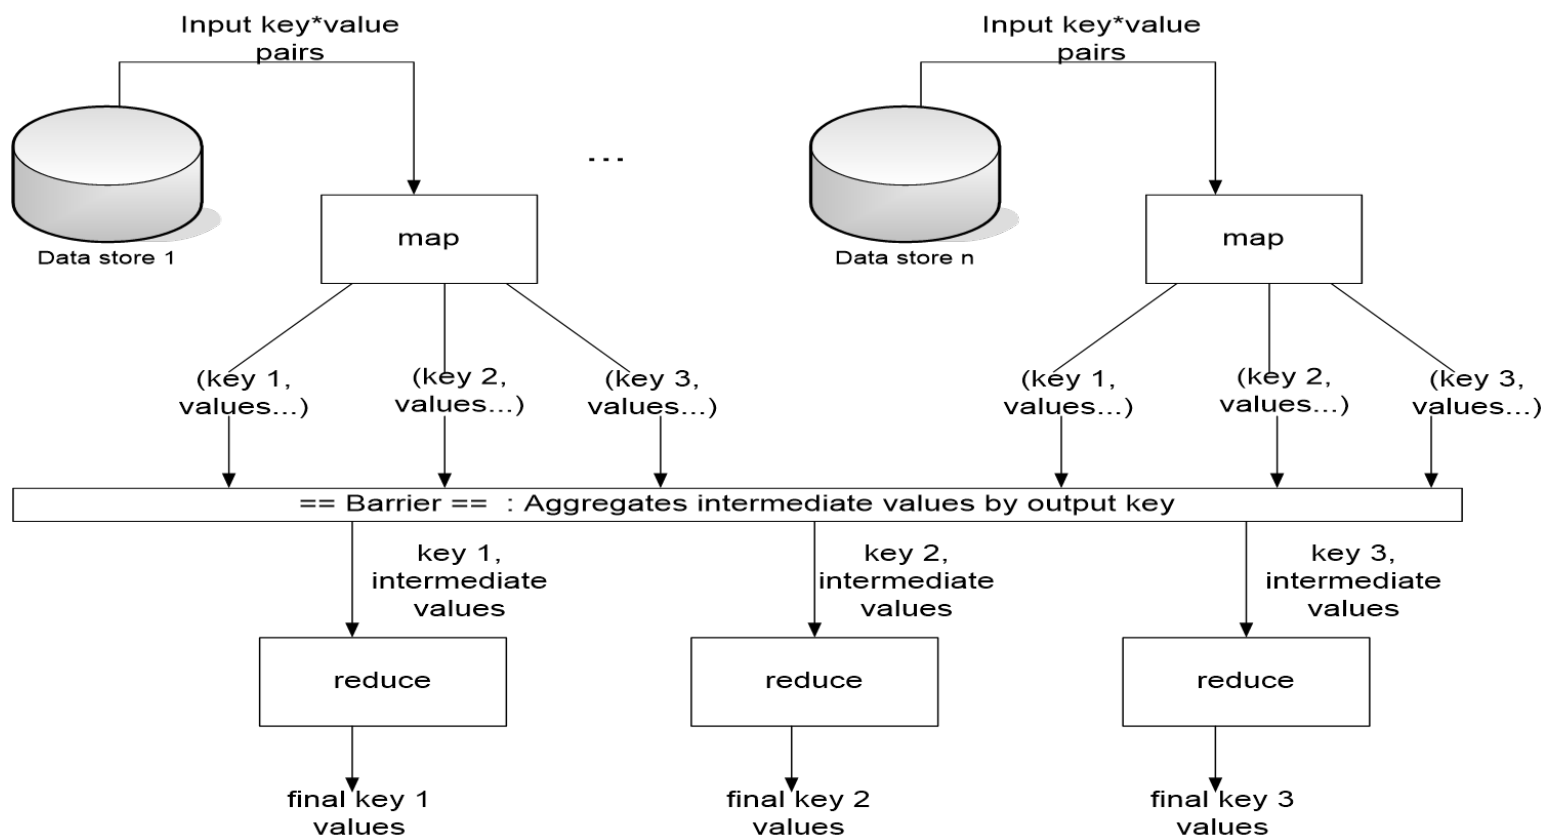
Records from the data source (lines out of files, rows of a database, etc) are fed into the map function as key\*value pairs: e.g., (filename, line).

map() produces one or more intermediate values along with an output key from the input.

After the map phase is over, all the intermediate values for a given output key are combined together into a list

reduce() combines those intermediate values into one or more final values for that same output key.(in practice, usually only one final value per key)

### 1.2.4 Architecture



### 1.3 简单的自我尝试

WordCount on Hadoop  
比较简单



## Chapter 2

# Mini Exercise

### 2.1 Exercise 1

#### 2.1.1 要求

Practise using basic hadoop command and fill in the following table

#### 2.1.2 实验过程

**Start hadoop**

**Use command to compute pi** <nMaps> is the number of mapper jobs and <nSamples> is the number of samples

```
hadoop jar /usr/local/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-  
examples-2.2.0.jar pi <nMaps> <  
nSamples>
```

Number of Maps	Number of samples	Time(s)	$\pi$
2	10	15.088 seconds	3.800
5	10	17.118	3.2800
10	10	18.179	3.200
2	100	14.996	3.1200
10	100	28.183	3.14800
100	100'0000	126.359 seconds	3.1415925600
100	1000'0000	133.294 seconds	3.14159273600
100	1'0000'0000	154.344 seconds	3.141592649200
1000	1'0000'0000	1417.978 seconds	3.1415926557200
2000	1'0000'0000	2628.074 seconds	3.1415926575600

随着测试数量的上升，时间和精确程度都在上升！

```

18/11/08 18:46:04 INFO mapreduce.Job: Job job_1541668829670_0009 completed successfully
18/11/08 18:46:04 INFO mapreduce.Job: Counters: 43
  File System Counters
    FILE: Number of bytes read=22006
    FILE: Number of bytes written=79680381
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=267890
    HDFS: Number of bytes written=215
    HDFS: Number of read operations=4003
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=3
  Job Counters
    Launched map tasks=1000
    Launched reduce tasks=1
    Data-local map tasks=1000
    Total time spent by all maps in occupied slots (ms)=5937503
    Total time spent by all reduces in occupied slots (ms)=1167148
  Map-Reduce Framework
    Map input records=1000
    Map output records=2000
    Map output bytes=18000
    Map output materialized bytes=28000
    Input split bytes=149890
    Combine input records=0
    Combine output records=0
    Reduce input groups=2
    Reduce shuffle bytes=28000
    Reduce input records=2000
    Reduce output records=0
    Spilled Records=4000
    Shuffled Maps =1000
    Failed Shuffles=0
    Merged Map outputs=1000
    GC time elapsed (ms)=72621
    CPU time spent (ms)=2787850
    Physical memory (bytes) snapshot=261899935744
    Virtual memory (bytes) snapshot=844577435648
    Total committed heap usage (bytes)=204668928000
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=118000
  File Output Format Counters
    Bytes Written=97
Job Finished in 1417.978 seconds
Estimated value of Pi is 3.14159265572000000000

```

展示 1000(Map) \* 1'0000'0000(Samples) 的结果!

Get the result

## 2.2 Exercise 2

### 2.2.1 要求

Work out a solution to make the computed pi approximate the 5th digit after the decimal dot correctly.

### 2.2.2 实验过程

实验过程和 1 类似, 为了进行更精确的计算, 调大了内存 (6G -> 7.5G), 分配了更多处理器内核 (6 -> 8), 观察是否有更快的结果输出。

```

18/11/08 19:50:53 INFO mapreduce.Job: Job job_1541674758306_0002 completed successfully
18/11/08 19:50:53 INFO mapreduce.Job: Counters: 43
  File System Counters
    FILE: Number of bytes read=44006
    FILE: Number of bytes written=159279379
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=534890
    HDFS: Number of bytes written=215
    HDFS: Number of read operations=8003
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=3
  Job Counters
    Launched map tasks=2000
    Launched reduce tasks=1
    Data-local map tasks=2000
    Total time spent by all maps in occupied slots (ms)=11088498
    Total time spent by all reduces in occupied slots (ms)=2182823
  Map-Reduce Framework
    Map input records=2000
    Map output records=4000
    Map output bytes=36000
    Map output materialized bytes=56000
    Input split bytes=298890
    Combine input records=0
    Combine output records=0
    Reduce input groups=2
    Reduce shuffle bytes=56000
    Reduce input records=4000
    Reduce output records=0
    Spilled Records=8000
    Shuffled Maps =2000
    Failed Shuffles=0
    Merged Map outputs=2000
    GC time elapsed (ms)=171213
    CPU time spent (ms)=5635570
    Physical memory (bytes) snapshot=531544735744
    Virtual memory (bytes) snapshot=1700306763776
    Total committed heap usage (bytes)=410098597888
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=236000
  File Output Format Counters
    Bytes Written=97
Job Finished in 2628.074 seconds
Estimated value of Pi is 3.14159265756000000000

```

展示 2000(Map) \* 1'0000'0000(Samples) 的结果!

结果是: 3.1415926575600 和 pi 的近似值 (3.1415926535898) 比较接近, 满足练习要求!

## Part II

# 实验总结

这次实验主要是学习新的 Hadoop，很有趣！

在配置环境的过程中，出现了一些奇妙的问题。通过使用 VMWare 的快照和自己 DeBUG 的过程学习到了新的 Ubuntu 的知识，有所进步！

期待在下一次实验中学习更多的知识！

## Part III

### 实验 9



原本已经写好了实验 8 的实验报告，然后上课前才得知实验 8 和实验 9 一起交，所以实验 9 的报告就接在实验 8 的报告后面啦。

## Chapter 3

# 实验准备

### 3.1 背景知识


#### 3.1.1 HDFS Concepts

**Very large** in this context means files that are hundreds of megabytes, gigabytes, or terabytes in size. There are Hadoop clusters running today that store petabytes of data.

**Streaming data access** HDFS is built around the idea that the most efficient data processing pattern is a write-once, read-many-times pattern.

**Commodity hardware** Hadoop doesn't require expensive, highly reliable hardware to run on.

### 3.1.2 Hadoop Streaming



## Chapter 4

# 实验操作

### 4.1 exercise1

实验一非常简单，按照要求写出 mapper.py 和 reducer.py 就可以了。

#### 4.1.1 mapper.py

```
import sys

# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1
        print '%s\t%s' % (word, 1)
```

#### 4.1.2 reducer.py

```
#!/usr/bin/env python

from operator import itemgetter
import sys

current_word = None
current_count = 0
current_len = 0
word = None

# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()

    # parse the input we got from mapper.py
    word, len = line.split('\t', 1)

    # convert count (currently a string) to int
    try:
        len = int(len)
    except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue

    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_len += len
        current_count += 1
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, round(current_len/current_count,2))
            current_len = len
            current_count=1
        current_word = word

    # do not forget to output the last word if needed!
    if current_word == word:
        print '%s\t%s' % (current_word, round(current_len/current_count,2))
```

在这次实验中还复习了 python 的相关操作，但总的来说比较基础就不再对代码进行过多的解释。

### 4.1.3 结果展示

```
hduser@ubuntu:~$ echo "we become what we do " |~/experiment/src2/ex1_mapper.py | sort -k1,1| ~/experiment/src2/ex1_reducer.py
tw      6.0
Sd      2.0
w       2.67
hduser@ubuntu:~$
```

使用新的测试语句进行测试

```
hduser@ubuntu:~$ echo "Portrait of Maria Portinari " |~/experiment/src2/ex1_mapper.py | sort -k1,1| ~/experiment/src2/ex1_reducer.py
M       5.0
o       2.0
P       8.5
hduser@ubuntu:~$
```

## 4.2 exercise2

这个实验应用的是迭代的操作计算 PageRank

我使用的方法和 accumulate 相似，不断迭代进行计算。

### 4.2.1 mapper.py

```
#!/usr/bin/env python

import sys

PL = {}
PN = {}

for line in sys.stdin:
    line = line.strip()
    if len(line) > 0:
        word, key = line.split('\t', 1)
        key = float(key)
        PL[word] = key

Ar = (0.15 / 4)
print 'A\t%f' % Ar
```

```

Br = (0.15 / 4) + 0.85 * ((PL['A'] / 3) + PL['D'])
print 'B\t%f' % Br
Cr = (0.15 / 4) + 0.85 * ((PL['A'] / 3) + (PL['B'] / 2))
print 'C\t%f' % Cr
Dr = (0.15 / 4) + 0.85 * ((PL['A'] / 3) + (PL['B'] / 2) + (PL['C']))
print 'D\t%f' % Dr

```

#### 4.2.2 结果展示

使用 python 直接运行 mapper.py

```

hduser@ubuntu:~/experiment/src/PageRank$ cat rank.txt | python mapper.py
A      0.037500
B      0.320833
C      0.214583
D      0.427083

```

使用 hadoop 的迭代方法，得到相同的结果。

```

hduser@ubuntu:~/experiment/src/PageRank$ ./batch_test.sh 1
Processing 1

```

A	0.037500
B	0.320833
C	0.214583
D	0.427083

结果

#### 4.2.3 在第一次完成实验的基础上，我又重新写了一份代码使用 mapper + reducer 的方法来进行操作

mapper.py 文件如下

```

#!/usr/bin/env python

import sys
#file1 = open("test2.txt")
# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into parts
    eles = line.split()

```

```

current_link = eles[0]
current_pr = float(eles[1])

print '%s\t%s\t%f' % (current_link, 0, 0.0375)

try:
    outlinks = eles[2:]

    outnum = float(len(outlinks))

    # increase counters
    for link in outlinks:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1

        print '%s\t%s\t%f' % (link, current_link, 0.85*current_pr/outnum)
except:
    pass

```

reducer.py 文件如下

```

#!/usr/bin/env python

from operator import itemgetter
import sys

current_page = None
current_pr = 0
dict1={}
dict2={}
for i in range(5):
    dict1[i] = []
    dict2[i] = 0;
#file1 = open("reduce.txt")
# input comes from STDIN
for line in sys.stdin:
    #print line
    line = line.strip()
    # parse the input we got from mapper.py
    try:
        page, target, pr = line.split()

    except:
        break

```



```

# convert count (currently a string) to int
try:
    pr = float(pr)
    target = int(target)
except ValueError:
    continue

dict1[target].append(page)
dict2[int(page)]+=pr

for key,value in dict2.items():
    if key != 0:
        print key,'\t',value,'\t','\t'.join(dict1[int(key)])

```

更新后的 batch\_test.sh 文件

```

#!/bin/bash

command='hadoop jar /usr/local/hadoop/share/hadoop/tools/lib/hadoop-
streaming-2.2.0.jar -files mapper.py,
reducer.py -mapper mapper.py -reducer
reducer.py'

mv='hadoop fs -mv '
rm='hadoop fs -rm -r '
cp2local='hadoop fs -copyToLocal '
input='tempinput'
for ((i=1;i<=1000;i++));
do
    echo "Processing $i"
    output="tempoutput_$i"
    eval "$command -input $input/* -output $output"
    input=$output
    eval "$rm $input/_SUCCESS"
done
mkdir /home/hduser/result
eval "$cp2local $output/* /home/hduser/result"

```

#### 4.2.4 实验结果展示

现在有了 mapper 和 reducer 之后，一切都变得顺理成章起来。使用语句让 hadoop 执行 1000 次之后，发现结果收敛。（其实不用那么多次，50-100 次过程中就可以看到明显的收敛结果）

```
hdsuser@ubuntu:~/Experiment/jars
Processing 1000
10/11/20 21:25:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
package30.jar: [/tmp/hadoop-huser/hadoop-unjar8705836796160083671/] [/tmp/streamjob5806197831362759388.jar tmpDir=null]
10/11/20 21:25:14 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
10/11/20 21:25:14 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
10/11/20 21:25:15 INFO mapred.FileInputFormat: Total input paths to process : 1
10/11/20 21:25:15 INFO mapreduce.JobSubmitter: number of splits:2
10/11/20 21:25:15 INFO Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.cache.files.filesizes is deprecated. Instead, use mapreduce.job.cache.files.filesizes
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.output.value.class is deprecated. Instead, use mapreduce.job.output.value.class
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.mapoutput.value.class is deprecated. Instead, use mapreduce.map.output.value.class
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.used.neweroptionsparser is deprecated. Instead, use mapreduce-client.neweroptionsparser.used
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.output.key.class is deprecated. Instead, use mapreduce.job.output.key.class
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.mapoutput.key.class is deprecated. Instead, use mapreduce.map.output.key.class
10/11/20 21:25:15 INFO Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
10/11/20 21:25:15 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1542699188409_1000
10/11/20 21:25:15 INFO hdfs.YarnClientImpl: Submitted application application_1542699188409_1000 to ResourceManager at /0.0.0.0:8032
10/11/20 21:25:15 INFO mapreduce.Job: The url to track the job: http://ubuntu:8088/proxy/application_1542699188409_1000/
10/11/20 21:25:15 INFO mapreduce.Job: Running job: job_1542699188409_1000
10/11/20 21:25:21 INFO mapreduce.Job: Job job_1542699188409_1000 running in uber mode : false
10/11/20 21:25:21 INFO mapreduce.Job: map 0% reduce 0%
10/11/20 21:25:26 INFO mapreduce.Job: map 100% reduce 0%
10/11/20 21:25:31 INFO mapreduce.Job: map 100% reduce 100%
10/11/20 21:25:31 INFO mapreduce.Job: Job job_1542699188409_1000 completed successfully
10/11/20 21:25:31 INFO mapreduce.Job: Counters: 43
File System Counters
  FILE: Number of Bytes read=171
  FILE: Number of Bytes written=24003
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of Bytes read=318
  HDFS: Number of Bytes written=64
  HDFS: Number of read operations=0
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=6514
  Total time spent by all reduces in occupied slots (ms)=2386
Map-Reduce Framework
  Map input records=4
  Map output records=11
```

```
hduser@ubuntu: ~/result
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:18 tempoutput_980
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:18 tempoutput_981
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:19 tempoutput_982
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:19 tempoutput_983
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:20 tempoutput_984
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:20 tempoutput_985
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:20 tempoutput_986
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:21 tempoutput_987
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:21 tempoutput_988
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:21 tempoutput_989
drwxr-xr-x - hduser supergroup 0 2018-11-20 16:10 tempoutput_99
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:22 tempoutput_990
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:22 tempoutput_991
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:22 tempoutput_992
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:23 tempoutput_993
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:23 tempoutput_994
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:23 tempoutput_995
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:24 tempoutput_996
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:24 tempoutput_997
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:24 tempoutput_998
drwxr-xr-x - hduser supergroup 0 2018-11-20 21:25 tempoutput_999
hduser@ubuntu:~$ hadoop fs -copyToLocal temp*txt ~/result
18/11/20 21:58:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyToLocal: 'temp*txt': No such file or directory
hduser@ubuntu:~$ cd ~/result/
hduser@ubuntu:~/result$ hadoop fs -ls tempoutput_999
18/11/20 21:59:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 64 2018-11-20 21:25 tempoutput_999/part-00000
hduser@ubuntu:~/result$ hadoop fs -copyToLocal tempoutput_999/part-00000.txt ~/result
18/11/20 22:01:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyToLocal: 'tempoutput_999/part-00000.txt': No such file or directory
hduser@ubuntu:~/result$ hadoop fs -copyToLocal tempoutput_999/part-00000.txt ~/result/result.txt
18/11/20 22:01:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyToLocal: 'tempoutput_999/part-00000.txt': No such file or directory
hduser@ubuntu:~/result$ hadoop fs -copyToLocal tempoutput_999/part-00000.txt ~/result/result.txt
18/11/20 22:02:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
copyToLocal: 'tempoutput_999/part-00000.txt': No such file or directory
hduser@ubuntu:~/result$ hadoop fs -cat tempoutput_999/part-00000
18/11/20 22:02:52 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platfo
rm... using builtin-java classes where applicable
1 0.0375 2 3 4
2 0.373247 3 4
3 0.206755 4
4 0.382497 2
hduser@ubuntu:~/result$
```

最后的结果：迭代 1 次：A 0.0375 B 0.320833 C 0.214583 D 0.429083

迭代 1000 次：A 0.0375 B 0.373247 C 0.206755 D 0.382497

## Part IV

# 实验总结

这次实验主要是学习新的 Hadoop 知识及新的应用，很有趣！

在实践过程中，由于第一次的实验深入了解了 hadoop 的知识，所以在 exercise 操作中写 py 文件都很顺利。但是也确实在学习过程中出现了很多基础语言不清楚、基础概念不明晰而导致的错误，所以会在基础的语言角度更多地再去理解！

期待在下一一次实验中学习更多的知识！